

BODIGE SAINATH

700756620

ICP – 7

Git – <https://github.com/BodigeSainath/icp7>

Video -

[https://drive.google.com/file/d/1Ko6Wpj1ls61\\_8usfNE1iWgawxbon6X3q/view?usp=sharing](https://drive.google.com/file/d/1Ko6Wpj1ls61_8usfNE1iWgawxbon6X3q/view?usp=sharing)

**1. Save the model and use the saved model to predict on new text data (ex, “A lot of good things are happening. We are respected again throughout the world, and that's a great thing.@realDonaldTrump”)**

```
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
from matplotlib import pyplot
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
import re

from sklearn.preprocessing import LabelEncoder

data = pd.read_csv('/content/Sentiment (3).csv')
# Keeping only the necessary columns
data = data[['text', 'sentiment']]

data['text'] = data['text'].apply(lambda x: x.lower())
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-z0-9\s]', '', x))

for idx, row in data.iterrows():
    row[0] = row[0].replace('rt', ' ')

max_fatures = 2000
tokenizer = Tokenizer(num_words=max_fatures, split=' ')
tokenizer.fit_on_texts(data['text'].values)
X = tokenizer.texts_to_sequences(data['text'].values)

X = pad_sequences(X)

embed_dim = 128
lstm_out = 196
```

```

def createmodel():
    model = Sequential()
    model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1]))
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
    model.add(Dense(3,activation='softmax'))
    model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics
= ['accuracy'])
    return model
# print(model.summary())

labelencoder = LabelEncoder()
integer_encoded = labelencoder.fit_transform(data['sentiment'])
y = to_categorical(integer_encoded)
X_train, X_test, Y_train, Y_test = train_test_split(X,y, test_size = 0.33,
random_state = 42)

batch_size = 32
model = createmodel()
model.fit(X_train, Y_train, epochs = 1, batch_size=batch_size, verbose = 2)
score,acc = model.evaluate(X_test,Y_test,verbose=2,batch_size=batch_size)
print(score)
print(acc)
print(model.metrics_names)

291/291 - 48s - loss: 0.8208 - accuracy: 0.6428 - 48s/epoch - 166ms/step
144/144 - 4s - loss: 0.7668 - accuracy: 0.6614 - 4s/epoch - 31ms/step
0.7668231725692749
0.6614242196083069
['loss', 'accuracy']

```

---

```

import tweepy
from keras.models import load_model
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import re

# Load the saved model
model = load_model("/content/sentiment_model.h5")

# Define a function for preprocessing text
def preprocess_text(text):
    text = text.lower()
    text = re.sub('[^a-zA-z0-9\s]', '', text)
    return text

# Example new text data

```

```
new_text = "A lot of good things are happening. We are respected again  
throughout the world, and that's a great thing. @realDonaldTrump"
```

```
# Preprocess the new text data
```

```
new_text = preprocess_text(new_text)
```

```
# Tokenize and pad the new text data
```

```
max_fatures = 2000
```

```
tokenizer = Tokenizer(num_words=max_fatures, split=' ')
```

```
tokenizer.fit_on_texts([new_text])
```

```
X_new = tokenizer.texts_to_sequences([new_text])
```

```
X_new = pad_sequences(X_new, maxlen=model.input_shape[1])
```

```
# Make predictions
```

```
predictions = model.predict(X_new)
```

```
# Determine the sentiment based on the prediction
```

```
sentiments = ['Negative', 'Neutral', 'Positive']
```

```
predicted_sentiment = sentiments[predictions.argmax()]
```

```
# Print the result
```

```
print("Predicted Sentiment: " + predicted_sentiment)
```

```
1/1 [=====] - 0s 296ms/step
```

```
Predicted Sentiment: Negative
```

---

## 2. Apply GridSearchCV on the source code provided in the class

```
!pip install scikeras
```

```
Collecting scikeras
```

```
  Downloading scikeras-0.12.0-py3-none-any.whl (27 kB)
```

```
Requirement already satisfied: packaging>=0.21 in  
/usr/local/lib/python3.10/dist-packages (from scikeras) (24.0)
```

```
Requirement already satisfied: scikit-learn>=1.0.0 in  
/usr/local/lib/python3.10/dist-packages (from scikeras) (1.2.2)
```

```
Requirement already satisfied: numpy>=1.17.3 in  
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.0.0->scikeras)  
(1.25.2)
```

```
Requirement already satisfied: scipy>=1.3.2 in  
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.0.0->scikeras)  
(1.11.4)
```

```
Requirement already satisfied: joblib>=1.1.1 in  
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.0.0->scikeras)  
(1.3.2)
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.0.0->scikeras)
(3.4.0)
```

```
Installing collected packages: scikeras
```

```
Successfully installed scikeras-0.12.0
```

```
from scikeras.wrappers import KerasClassifier
```

```
import pandas as pd
```

```
import re
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
```

```
from tensorflow.keras.utils import to_categorical
```

```
from sklearn.model_selection import train_test_split, GridSearchCV
```

```
from sklearn.preprocessing import LabelEncoder
```

```
from scikeras.wrappers import KerasClassifier
```

```
# Assuming the data loading and preprocessing steps are the same
```

```
max_features = 2000
```

```
tokenizer = Tokenizer(num_words=max_features, split=' ')
```

```
# Assuming tokenizer fitting and text preprocessing is done here
```

```
def createmodel(optimizer='adam'):
```

```
    model = Sequential()
```

```
    model.add(Embedding(max_features, embed_dim, input_length=X.shape[1]))
```

```
    model.add(SpatialDropout1D(0.2))
```

```
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
```

```
    model.add(Dense(3, activation='softmax'))
```

```
    model.compile(loss='categorical_crossentropy', optimizer=optimizer,
metrics=['accuracy'])
```

```
    return model
```

```
# Define the KerasClassifier with the build_fn as our model creation function
```

```
model = KerasClassifier(model=createmodel, verbose=2)
```

```
# Define hyperparameters to tune
```

```
param_grid = {
```

```
    'batch_size': [32, 64],
```

```
    'epochs': [1, 2],
```

```
    'optimizer': ['adam', 'rmsprop']
```

```
}
```

```
# Initialize GridSearchCV
```

```
grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=1, cv=3)
```

```
# Fit GridSearchCV
```

```
grid_result = grid.fit(X_train, Y_train)
```

```
# Summarize results
```

```
print("Best: %f using %s" % (grid_result.best_score_,  
grid_result.best_params_))
```

```
194/194 - 37s - loss: 0.8596 - accuracy: 0.6328 - 37s/epoch - 192ms/step  
97/97 - 2s - 2s/epoch - 23ms/step  
194/194 - 41s - loss: 0.8563 - accuracy: 0.6297 - 41s/epoch - 210ms/step  
97/97 - 3s - 3s/epoch - 34ms/step  
194/194 - 36s - loss: 0.8773 - accuracy: 0.6278 - 36s/epoch - 186ms/step  
97/97 - 2s - 2s/epoch - 23ms/step  
194/194 - 32s - loss: 0.8712 - accuracy: 0.6326 - 32s/epoch - 167ms/step  
97/97 - 3s - 3s/epoch - 28ms/step  
194/194 - 33s - loss: 0.8588 - accuracy: 0.6292 - 33s/epoch - 171ms/step  
97/97 - 3s - 3s/epoch - 27ms/step  
194/194 - 34s - loss: 0.8675 - accuracy: 0.6252 - 34s/epoch - 173ms/step  
97/97 - 2s - 2s/epoch - 23ms/step  
Epoch 1/2  
194/194 - 33s - loss: 0.8632 - accuracy: 0.6300 - 33s/epoch - 171ms/step  
Epoch 2/2  
194/194 - 29s - loss: 0.7171 - accuracy: 0.6888 - 29s/epoch - 151ms/step  
97/97 - 3s - 3s/epoch - 32ms/step  
Epoch 1/2  
194/194 - 33s - loss: 0.8599 - accuracy: 0.6271 - 33s/epoch - 170ms/step  
Epoch 2/2  
194/194 - 30s - loss: 0.6978 - accuracy: 0.6991 - 30s/epoch - 157ms/step  
97/97 - 2s - 2s/epoch - 23ms/step  
Epoch 1/2  
194/194 - 35s - loss: 0.8553 - accuracy: 0.6285 - 35s/epoch - 179ms/step  
Epoch 2/2  
194/194 - 29s - loss: 0.6883 - accuracy: 0.7022 - 29s/epoch - 151ms/step  
97/97 - 2s - 2s/epoch - 23ms/step  
Epoch 1/2  
194/194 - 35s - loss: 0.8565 - accuracy: 0.6320 - 35s/epoch - 178ms/step  
Epoch 2/2  
194/194 - 29s - loss: 0.7122 - accuracy: 0.6949 - 29s/epoch - 150ms/step  
97/97 - 3s - 3s/epoch - 34ms/step  
Epoch 1/2  
194/194 - 33s - loss: 0.8660 - accuracy: 0.6295 - 33s/epoch - 168ms/step  
Epoch 2/2  
194/194 - 30s - loss: 0.7025 - accuracy: 0.6999 - 30s/epoch - 157ms/step  
97/97 - 2s - 2s/epoch - 23ms/step  
Epoch 1/2  
194/194 - 35s - loss: 0.8494 - accuracy: 0.6320 - 35s/epoch - 181ms/step  
Epoch 2/2  
194/194 - 30s - loss: 0.6845 - accuracy: 0.7093 - 30s/epoch - 156ms/step  
97/97 - 3s - 3s/epoch - 30ms/step  
97/97 - 30s - loss: 0.8820 - accuracy: 0.6182 - 30s/epoch - 309ms/step  
49/49 - 3s - 3s/epoch - 51ms/step  
97/97 - 28s - loss: 0.8731 - accuracy: 0.6228 - 28s/epoch - 290ms/step  
49/49 - 3s - 3s/epoch - 52ms/step  
97/97 - 30s - loss: 0.8955 - accuracy: 0.6165 - 30s/epoch - 307ms/step
```

49/49 - 2s - 2s/epoch - 51ms/step  
97/97 - 29s - loss: 0.8696 - accuracy: 0.6263 - 29s/epoch - 298ms/step  
49/49 - 2s - 2s/epoch - 50ms/step  
97/97 - 29s - loss: 0.8740 - accuracy: 0.6218 - 29s/epoch - 304ms/step  
49/49 - 3s - 3s/epoch - 65ms/step  
97/97 - 28s - loss: 0.8783 - accuracy: 0.6241 - 28s/epoch - 289ms/step  
49/49 - 3s - 3s/epoch - 67ms/step  
Epoch 1/2  
97/97 - 29s - loss: 0.8779 - accuracy: 0.6242 - 29s/epoch - 302ms/step  
Epoch 2/2  
97/97 - 25s - loss: 0.7220 - accuracy: 0.6949 - 25s/epoch - 259ms/step  
49/49 - 3s - 3s/epoch - 68ms/step  
Epoch 1/2  
97/97 - 29s - loss: 0.8862 - accuracy: 0.6176 - 29s/epoch - 303ms/step  
Epoch 2/2  
97/97 - 25s - loss: 0.7242 - accuracy: 0.6894 - 25s/epoch - 254ms/step  
49/49 - 2s - 2s/epoch - 50ms/step  
Epoch 1/2  
97/97 - 28s - loss: 0.8839 - accuracy: 0.6164 - 28s/epoch - 287ms/step  
Epoch 2/2  
97/97 - 25s - loss: 0.7149 - accuracy: 0.6877 - 25s/epoch - 255ms/step  
49/49 - 3s - 3s/epoch - 52ms/step  
Epoch 1/2  
97/97 - 30s - loss: 0.8833 - accuracy: 0.6216 - 30s/epoch - 309ms/step  
Epoch 2/2  
97/97 - 26s - loss: 0.7304 - accuracy: 0.6931 - 26s/epoch - 272ms/step  
49/49 - 4s - 4s/epoch - 83ms/step  
Epoch 1/2  
97/97 - 39s - loss: 0.8786 - accuracy: 0.6179 - 39s/epoch - 398ms/step  
Epoch 2/2  
97/97 - 27s - loss: 0.7233 - accuracy: 0.6889 - 27s/epoch - 278ms/step  
49/49 - 4s - 4s/epoch - 83ms/step  
Epoch 1/2  
97/97 - 33s - loss: 0.8707 - accuracy: 0.6198 - 33s/epoch - 336ms/step  
Epoch 2/2  
97/97 - 30s - loss: 0.7207 - accuracy: 0.6833 - 30s/epoch - 308ms/step  
49/49 - 3s - 3s/epoch - 52ms/step  
Epoch 1/2  
291/291 - 49s - loss: 0.8301 - accuracy: 0.6416 - 49s/epoch - 170ms/step  
Epoch 2/2  
291/291 - 46s - loss: 0.6884 - accuracy: 0.7066 - 46s/epoch - 158ms/step  
Best: 0.672548 using {'batch\_size': 32, 'epochs': 2, 'optimizer': 'adam'}